=== VERSION 0: NovaCred Flex UW Manual ===

NovaCred Financial

NovaCred Flex Installment Loan Underwriting Policy Manual

VERSION 0 | SYNTHETIC TRAINING ARTIFACT

This document is completely fictional. All names, rule names, thresholds, codes, decisions, and operating logic are invented for testing and retrieval evaluation.

No real lender policy, legal threshold, or production rule is represented.

PRODUCT OVERVIEW

The NovaCred Flex Installment Loan is a short-term consumer installment product intended to provide predictable, fixed-payment access to a modest line of credit for applicants who satisfy NovaCred's fictional identity, credit, income, fraud, compliance, and affordability controls. The product is designed around a standardized decision engine that can receive prescreened leads or evaluate a fresh application, apply deterministic rules in a documented sequence, and produce an underwriting outcome that can be explained to operations teams. The product may be offered through approved digital channels, relationship marketing, or controlled partner placements. The loan amount, term, pricing, and servicing treatment are determined only after the applicable underwriting path completes its required checks and the decision result is assigned. All examples in this manual are synthetic and are intentionally constructed to create varied retrieval and exception patterns for testing a downstream RAG system.

NovaCred Financial states that it seeks to administer the NovaCred Flex Installment Loan fairly and consistently and does not intentionally make credit decisions on the basis of protected characteristics or any other characteristic prohibited by applicable law. This fictional statement is paired with an internal expectation that the same data definitions, versioned rule logic, and exception-handling standards are applied to similarly situated applicants. The underwriting engine is a configurable decision service rather than a free-form judgment tool: where the manual specifies an automatic action, the action is expected to be taken unless an explicitly documented override or system exception applies. Manual Review is reserved for cases where the policy requires human assessment, evidence reconciliation, or a fraud/compliance investigation that cannot be responsibly resolved by deterministic automation.

MARKETING CHANNELS AND UNDERWRITING STRATEGIES

NovaCred supports two marketing and underwriting paths for the NovaCred Flex Installment Loan: Prescreen and Non-Prescreen. The Prescreen path begins with a controlled audience assembled from permitted data sources and is intended to support a firm or near-firm invitation after a preliminary eligibility screen. The Non-Prescreen path begins with an applicant who has not been selected through the Prescreen audience process. Both paths ultimately use the same conceptual rule engine, but they may apply different parameter values, sequence gates, data freshness windows, and verification requirements. A path identifier must travel with the application so downstream reviewers can reconstruct which population logic was used.

The two-path structure should not be interpreted to mean that the Non-Prescreen path is inherently weaker or stronger. Rather, the paths are designed around different information states. Prescreen applicants may have already satisfied an initial audience-selection screen, allowing the full underwriting stage to focus more heavily on application consistency and refreshed evidence. Non-Prescreen applicants generally require a broader discovery step before final decisioning. Any path-specific rule must be tagged as path-specific in the implementation configuration; identical rules should retain identical codes so that monitoring, issue tracking, and model-risk reporting can compare the paths without an artificial code mismatch.

PRESCREEN SELECTIONS

The fictional Prescreen audience is generated from an approved monthly selection file and a documented eligibility filter. The selection process is designed to reduce obvious ineligible records before an offer is sent while retaining enough volume for downstream application conversion. A candidate must satisfy all mandatory criteria unless the source-system status explicitly permits a documented exception. Prescreen selection criteria are intentionally narrower than the full underwriting rules: they are designed to identify candidates who can plausibly reach underwriting, not to replace underwriting itself.

Table 1: Prescreen Criteria

Criterion	Requirement	Source
Residence Status	U.S. resident status coded Eligible and no unresolved residency discrepancy	CivicGraph Residence File
Age Band	Applicant age at least 20 years and no age conflict across source feeds	NovaCred Identity Hub
Existing Exposure	No active Flex loan with a remaining balance above the synthetic exposure ceiling of $6,400	NovaCred Account Ledger
Recent Delinquency	No account-level delinquency status of 60+ days during the prior 180-day window	NovaCred Account Ledger

Criterion	Requirement	Source
Credit Availability	No active bureau freeze indicator in the monthly prescreen extract	OrbitScore Bureau Feed
Fraud Disposition	No unresolved hard-block fraud flag in the prior 12 months	NovaShield Fraud File
Contactability	Phone and email contact points each pass basic deliverability checks	Contact Integrity Service
Marketing Suppression	Record is not present on the internal do-not-market suppression list	NovaCred Suppression Register
Income Proxy	Synthetic income proxy at or above $1,900 monthly and not missing	NovaCred Affordability Index
Application Velocity	No more than three NovaCred-originated applications in the prior 45 days	NovaCred Application Ledger

NON-PRESCREEN MARKETING

Non-Prescreen marketing may be initiated from a general product landing page, a partner campaign, a controlled retargeting list, or another approved channel that does not originate from the Prescreen selection file. The applicant enters underwriting without relying on the prior audience-selection decision. The same rule engine is used, but path parameters may be more conservative in places where information is absent, stale, or less independently verified. For example, a Non-Prescreen applicant may be required to supply additional income evidence even when the Prescreen candidate previously passed a source-based income proxy filter. This difference is intentional and should not be implemented as separate rule definitions when the underlying rule objective is identical.

Marketing operations must preserve the campaign or source identifier, first-contact timestamp, and path flag. These fields are not themselves credit-decision variables unless explicitly listed in the underwriting configuration. They are retained to support monitoring of funnel behavior, offer aging, source-level conversion, and post-decision quality. A campaign record that is missing the path flag should be quarantined rather than silently assigned to Prescreen or Non-Prescreen based on an arbitrary default.

APPLICANT TYPES

Table 2: Applicant Types

Type	Description	Notes
Primary	The person whose application establishes the principal underwriting record and repayment obligation.	Required for all funded accounts.
Co-applicant	A second applicant evaluated alongside the Primary with a separate identity and income profile.	Not all channels permit co-applicants; both records must pass identity controls.
Joint	A relationship structure in which two applicants share the contractual borrowing record under the same decision package.	Joint applicants are evaluated together but each retains individual verification attributes.
Authorized Payer	A person allowed to provide a payment account or payment instruction without becoming a borrower.	Does not independently satisfy borrower underwriting requirements.
Business-Related Natural Person	A consumer applicant whose stated repayment source is partly derived from a sole-proprietor or informal business activity.	Subject to self-employed income documentation rules where applicable.

Applicant type must be established before the final rule sequence begins because several checks are keyed to the number of obligors. A Primary record with a co-applicant should not be treated as a single-obligor case for income sufficiency, identity matching, or document completeness. Conversely, an Authorized Payer should never be treated as an additional borrower merely because the payer account is attached to the application. The application orchestration layer must carry an explicit relationship code rather than attempting to infer type from free-text comments.

UNDERWRITING

Underwriting is organized into prequalification and full underwriting stages. Prequalification establishes whether the application can continue through the defined path using early identity, bureau, fraud, and basic affordability signals. Full underwriting incorporates refreshed evidence, final compliance checks, scoring, verification outcomes, and any required manual review. The rule engine should evaluate rules in their configured sequence; however, a later-stage rule may be skipped when an earlier hard disposition makes the remaining steps unnecessary. The skip reason must be captured so that monitoring can distinguish "not applicable" from "not evaluated due to terminal decision."

The rules below use invented three-digit codes. Codes are stable identifiers, not rank scores. Rule tables identify the control, source, action, and code; the detailed meaning and trigger definition for each code is maintained in Exhibit C so a code lookup can be answered by retrieving the rule row together with the applicable code-definition entry. A higher code number does not imply more severe outcome, later sequence position, or stronger control.

PRESCREEN UNDERWRITING RULES

The Prescreen path begins with a lightweight prequalification pass. Rules in Table 3 are intended to detect hard disqualifiers and obvious data deficiencies before deeper evidence is requested. In a production implementation, the sequence number would be stored separately from the display table; this manual shows the policy concept rather than a deployment schema.

Table 3: Prescreen Prequalification Rules

Rule	Source	Action	Code
Blocked SSN[1]	NovaShield Identity Blocklist	NOAA	101
Blocked Phone	NovaShield Contact Risk Feed	NOAA	102
Blocked Email	NovaShield Contact Risk Feed	NOAA	103
IP Velocity	NovaShield Network Monitor	Manual Review	104
Credit Freeze	OrbitScore Bureau Feed	RFAI	105
Thin File	OrbitScore Bureau Feed	RFAI	106
Score Floor	OrbitScore Bureau Feed	NOAA	107
Insufficient Income	NovaCred Affordability Index	NOAA	108
SSI Threshold	Income Verification Service	NOAA	109
Recent Decline	NovaCred Application Ledger	NOAA	110
Inquiry Velocity	OrbitScore Bureau Feed	NOAA	111
MLA Check	NovaCred Compliance Service	Manual Review	112
OFAC Match	NovaCred Compliance Service	NOAA	113
FACTA Alert	OrbitScore Bureau Feed	RFAI	114
Scoring	NovaCred Model Services	RFAI	115,116
Existing Exposure	NovaCred Account Ledger	NOAA	117
Duplicate Application	NovaCred Application Ledger	NOAA	118

Rule	Source	Action	Code
Contact Inconsistency	NovaCred Identity Hub	Manual Review	119

[1] Blocked SSN exception: if the block record is older than 36 months, is categorized as a resolved administrative duplicate, and the current identity comparison is exact, the result is overridden and treated as passed.

[2] Blocked Routing exception: if the prior block was generated solely from a temporary network anomaly and the current application passes the documented Intraday Account Review criteria, the rule result may be overridden and treated as passed.

Table 3 illustrates a deliberate distinction between hard outcomes and information requests. A blocked identity record normally generates a NOAA because continued processing is not permitted under the fictional control. A credit freeze or thin-file condition instead generates an RFAI because the policy contemplates additional evidence or a refreshed bureau state. IP Velocity is routed to Manual Review because velocity can be a legitimate artifact of shared networks, partner traffic, or repeated applicant behavior and therefore requires contextual assessment.

Full underwriting for Prescreen candidates does not simply repeat the first-stage checks. It refreshes volatile attributes, verifies the application against current evidence, and applies final affordability and compliance controls. Identical controls reuse codes, while materially different final-stage logic receives a new code. Reviewers should not infer that a code reused across stages is redundant: the source event, timestamp, and evaluated state can differ even when the policy objective is the same.

Table 4: Prescreen Full Underwriting Rules

Rule	Source	Action	Code
Blocked Routing[2]	NovaShield Identity Blocklist	NOAA	120
Credit Score Confirmation	OrbitScore Bureau Feed	NOAA	121
Credit Freeze Refresh	OrbitScore Bureau Feed	RFAI	105
Thin File Recheck	OrbitScore Bureau Feed	RFAI	106
Income Reconciliation	Income Verification Service	RFAI	122
Employment Tenure	Income Verification Service	NOAA	123
Bank Stability	NovaCred Transaction Analyzer	NOAA	124
Payment Stress	NovaCred Affordability Index	NOAA	125
MLA Final	NovaCred Compliance Service	Manual Review	112
OFAC Final	NovaCred Compliance Service	NOAA	113
FACTA Resolution	OrbitScore Bureau Feed	RFAI	114
Scoring[3]	NovaCred Model Services	RFAI	115,116
High Inquiry Overlay	OrbitScore Bureau Feed	NOAA	126
Recent Application Decline	NovaCred Application Ledger	NOAA	110
Manual Identity Reconciliation	NovaCred Identity Hub	Manual Review	127
Offer Age	NovaCred Marketing Ledger	RFAI	128
Residence Mismatch	NovaCred Identity Hub	Manual Review	129

Rule	Source	Action	Code
Document Integrity	NovaCred Document Gateway	NOAA	130

[3] Scoring exception: if the credit score vector is unavailable and the Null Score condition did not trigger, the scoring result is overridden and treated as passed for the purpose of continuing to non-score controls. The override does not create a score and does not guarantee approval.

When a rule produces RFAI, the applicant record enters an information-request state rather than an automatic decline. The requested evidence must be mapped to a specific reason code and may have an expiration date. When a rule produces Manual Review, operations should not convert the result to RFAI merely to avoid human intervention; the manual-review condition exists because the policy anticipates a judgment or reconciliation task. Conversely, manual reviewers may not invent alternate decline reasons that are unsupported by the configured policy.

NON-PRESCREEN UNDERWRITING RULES

The Non-Prescreen path uses the same core identity, bureau, income, compliance, and scoring controls but can be more conservative when source evidence is incomplete or the lead did not originate from the controlled Prescreen population. This section intentionally reuses codes for identical rules so that aggregate monitoring can identify whether the same control behaves differently by path. Three additional controls are unique to Non-Prescreen because they address the higher uncertainty of direct acquisition traffic.

Table 5: Non-Prescreen Prequalification Rules

Rule	Source	Action	Code
Blocked SSN	NovaShield Identity Blocklist	NOAA	101
Blocked Phone	NovaShield Contact Risk Feed	NOAA	102
Blocked Email	NovaShield Contact Risk Feed	NOAA	103
IP Velocity	NovaShield Network Monitor	Manual Review	104
Credit Freeze	OrbitScore Bureau Feed	RFAI	105
Thin File	OrbitScore Bureau Feed	RFAI	106
Score Floor	OrbitScore Bureau Feed	NOAA	107
Insufficient Income	NovaCred Affordability Index	NOAA	108
SSI Threshold	Income Verification Service	NOAA	109
MLA Check	NovaCred Compliance Service	Manual Review	112
OFAC Match	NovaCred Compliance Service	NOAA	113
FACTA Alert	OrbitScore Bureau Feed	RFAI	114
Scoring	NovaCred Model Services	RFAI	115,116
Acquisition Source Risk	NovaCred Channel Monitor	Manual Review	131
Device Reuse	NovaShield Device Graph	Manual Review	132
Early-Payoff Pattern	NovaCred Account Ledger	NOAA	133

[1] Non-Prescreen blocked SSN exception follows the same exception defined for Table 3.

[4] See Table 10 for the full list of applicable RFAI codes used by the information-request workflow.

Table 6: Non-Prescreen Full Underwriting Rules

Rule	Source	Action	Code
Blocked Routing	NovaShield Identity Blocklist	NOAA	120
Credit Score Confirmation	OrbitScore Bureau Feed	NOAA	121
Credit Freeze Refresh	OrbitScore Bureau Feed	RFAI	105
Thin File Recheck	OrbitScore Bureau Feed	RFAI	106
Income Reconciliation	Income Verification Service	RFAI	122
Employment Tenure	Income Verification Service	NOAA	123
Bank Stability	NovaCred Transaction Analyzer	NOAA	124
Payment Stress	NovaCred Affordability Index	NOAA	125
MLA Final	NovaCred Compliance Service	Manual Review	112
OFAC Final	NovaCred Compliance Service	NOAA	113
FACTA Resolution	OrbitScore Bureau Feed	RFAI	114
Scoring	NovaCred Model Services	RFAI	115,116
High Inquiry Overlay	OrbitScore Bureau Feed	NOAA	126
Manual Identity Reconciliation	NovaCred Identity Hub	Manual Review	127
Offer Age	NovaCred Marketing Ledger	RFAI	128
Residence Mismatch	NovaCred Identity Hub	Manual Review	129
Document Integrity	NovaCred Document Gateway	NOAA	130
Channel Attestation	NovaCred Channel Monitor	RFAI	134
Address Evidence	NovaCred Identity Hub	RFAI	139
Funding Method Verification	NovaCred Payment Graph	RFAI	140
Bank Account Ownership	NovaCred Payment Graph	NOAA	135
Notice Generation	NovaCred Notice Service	NOIA	136
Data Correction Notice	NovaCred Notice Service	NOIA	137
Disposition Notice Hold	NovaCred Notice Service	NOIA	138

Path-specific differences should be visible in configuration rather than hidden in code. For example, the Non-Prescreen path adds Acquisition Source Risk, Device Reuse, Early-Payoff Pattern, Channel Attestation, and Bank Account Ownership because direct or partner-acquired traffic can produce evidence patterns not present in the controlled Prescreen file. Those rules have dedicated codes. A common mistake in implementation is to clone the entire rule table and assign new codes to every row; that makes model monitoring and policy comparison harder and obscures whether two paths are truly using the same control.

SCORE CUT AND LINE ASSIGNMENT

Line assignment is a two-axis lookup using a refreshed Vantage Score band and the NovaCred Risk Score band. The adjustment shown below is applied to a synthetic base line after all hard policy rules have passed or been appropriately resolved. The matrix deliberately increases the adjustment as both score dimensions improve. The matrix is not itself an approval rule: a positive adjustment cannot override a decline, compliance stop, or unresolved manual-review condition. A negative value reduces the proposed line, and a zero value retains the applicable base line.

Table 7: Line Assignment Matrix

Vantage Score band	<680	680-749	750-829	830-909	>909
<530	-$1,500	-$1,250	-$1,000	-$750	-$500
530-569	-$1,250	-$900	-$650	-$400	-$200
570-609	-$900	-$600	-$250	+$100	+$250
610-649	-$500	-$200	+$150	+$450	+$700
>649	-$250	+$150	+$500	+$850	+$1,100

[1] Minimum line floor applies regardless of adjustment.

[2] Maximum line cap applies regardless of adjustment.

The matrix should be evaluated only after both score bands are normalized to the current scoring schema. A missing Vantage band or a missing NovaCred Risk Score band should not be converted to the highest or lowest band by default. The scoring override in Table 4 addresses one very specific missing-vector scenario and does not authorize line assignment from an invented or stale score. Where the score cannot be normalized to a valid band, operations must use the applicable information-request or manual-review path.

FRAUD, VERIFICATIONS, AND MANUAL REVIEW

NovaCred's fictional fraud process is designed as a layered control rather than a single fraud score. Automated identity, contact, network, device, and account-history signals are evaluated first. When evidence is strongly adverse, the engine can issue a NOAA and stop the application. When the evidence is ambiguous but potentially curable, the engine requests additional information. When the evidence requires contextual comparison or source reconciliation, the application is routed to Manual Review. Reviewers must document the evidence considered, the disposition selected, and any exception authority used. Manual review is not an unrestricted override channel; it is an explicit control intended to resolve cases that the automated rule set cannot safely classify.

A fraud reviewer should compare signals across time rather than treating one isolated indicator as conclusive. For example, a device reused across multiple applicants may be suspicious, but the same pattern can arise in a shared household or managed-device environment. Similarly, a contact point that matches several applications may reflect synthetic coordination, but may also be present on a common business or family account. The policy therefore distinguishes high-confidence blocks from contextual triggers and requires a documented reason when a Manual Review case is returned to the automated flow.

INTRADAY ACCOUNT REVIEW

Intraday Account Review is a rapid control used when activity changes materially after the initial application decision but before funding or first disbursement. It evaluates newly observed application velocity, device reuse, contact changes, payment-account linkage, and other fast-moving signals. The review should be performed using the latest available event snapshot and should not silently replace the original underwriting record. When an intraday review reverses a previously favorable state, the system should preserve both the original decision and the subsequent reversal reason for auditability.

Fraud Consortium criteria are detailed in Table 14. The consortium is a fictional shared-signal layer that groups synthetic patterns such as repeated device identity, high-frequency contact reuse, and coordinated account behavior. A consortium match is not necessarily an automatic decline; the disposition depends on the specific criterion, confidence level, and whether the application has independently cleared other identity controls.

INCOME VERIFICATION

Income verification establishes whether the repayment-capacity assumptions used by the underwriting engine are supported by sufficiently reliable evidence. The product uses a tiered approach so that lower-risk applications can be verified with lower-friction evidence while cases with more variable or complex income require stronger documentation. Income should be normalized into a monthly amount using the designated synthetic convention for the source type. Reviewers must not count the same income stream twice merely because it appears in multiple documents or transaction feeds.

Table 8: Income Verification Tiers

Tier	Criteria	Required Documents
Tier 1	Stable salaried employment; automated source is current and internally consistent.	One current pay statement plus automated employment confirmation.
Tier 2	Variable employment, recent job change, or incomplete automated confirmation.	Two recent pay statements and one employment-status confirmation.
Tier 3	Self-employed, contract-based, commission-heavy, or business-derived income.	Two recent bank statements, one business/income statement, and the applicable self-employment attestation.[1]
Tier 4	Complex or disputed income with material reconciliation variance.	Document package specified by Manual Review, plus source reconciliation notes and reviewer disposition.

[1] Tier 3 applies different rules for self-employed applicants: verified business inflow may be normalized after eligible expense adjustments and must not be treated as equivalent to gross receipts without the prescribed reconciliation.

[2] A document request generated by an RFAI code remains active only for the configured evidence window; submission after expiration may require a refreshed decision.

Income evidence is considered sufficient only when the data source, date, applicant relationship, and normalized amount can be established. A document that displays a large annual figure but lacks a reasonable time period should not be converted directly into a monthly income value. Likewise, a bank transaction stream that includes internal transfers must be normalized so that movement between the applicant's own accounts does not inflate repayment capacity. Where source systems disagree, the reviewer should identify the controlling source according to the tier rules rather than selecting the most favorable figure.

ACCOUNT MANAGEMENT

Post-origination account management governs selected events that occur after funding, including payment status changes, repeated failed debit attempts, material contact changes, and reapplication behavior. The account-management layer is separate from initial underwriting but may feed future underwriting through the Account Ledger and fraud graph. Post-origination monitoring does not retroactively change the original underwriting decision; instead, it creates a new event record that can influence a later application or servicing intervention.

The fictional policy expects operations teams to preserve a clean boundary between underwriting evidence and servicing evidence. A payment failure after funding should not be rewritten as evidence that the original income rule failed unless a later investigation establishes that the underwriting record itself was materially inaccurate. Similarly, a customer who repays early should not be labeled as higher risk solely because of early payoff; the Early-Payoff Pattern rule in Table 5 is designed only to detect repeated rapid payoff-and-reapply cycles that may indicate a different behavior pattern.

CONTROL INTERPRETATION AND DATA FRESHNESS

Freshness control: use the most recent approved source snapshot available for the underwriting stage, and retain its retrieval timestamp beside the normalized value used by the rule.

Stage-specific evidence should be immutable once a decision event is finalized. A later source refresh belongs to a new evaluation event rather than silently replacing the original evidence.

Source precedence must be explicit when two providers return conflicting values. The configuration should identify the preferred source, the permitted fallback, and the condition that authorizes the fallback.

Null and stale values should remain distinguishable. A missing field can trigger an information request, while an expired snapshot may require a refresh even when the field itself is present.

Review records should preserve both the raw status returned by the provider and the interpreted status used by the rules engine. This supports later reconstruction of mapping errors.

When a rule is rerun after an evidence update, the system should record which source change caused the rerun. Analysts should be able to distinguish a policy change from a data refresh.

Operational monitoring should flag unusual changes in source freshness, response completeness, or fallback usage because these conditions can alter rule-firing rates without any policy edit.

Manual reviewers must use the evidence version associated with the active evaluation event. A document uploaded later should be evaluated under the rework or refresh procedure rather than attached retroactively to an older decision.

MANUAL REVIEW PRACTICE

Manual Review begins with a coherence check across the application, identity evidence, transaction information, and bureau response. The reviewer should identify contradictions before deciding whether any condition can be cleared.

The reviewer should separate evidence collection from outcome selection. Evidence is recorded first; the applicable policy disposition is selected second. This reduces the risk of searching only for facts that support a preferred outcome.

Numeric rule codes are identifiers, not severity rankings. A reviewer should use the rule name, description, action, and applicable footnote to interpret a result rather than infer meaning from the code number.

A second-level review is warranted when the evidence clears one trigger but introduces an unresolved conflict elsewhere in the file. The escalation record should identify both the cleared condition and the remaining question.

For identity concerns, compare the submitted value with the source response and record the exact mismatch. Do not summarize a mismatch merely as 'identity issue' when the underlying field is known.

For income concerns, document the income period, source type, normalization method, and any expense adjustment used in the decision. Annualized values should not be substituted without recording the transformation.

For bureau exceptions, identify whether the issue is a freeze, thin file, missing score, or stale response. These conditions may lead to different information or review paths even when they all interrupt scoring.

When an automated fraud signal is challenged by contextual evidence, the reviewer should preserve both sides of the record. The exception should state what new fact changed the assessment and who authorized the resolution.

Manual Review should not erase the original automated event. The final record should contain the triggered rule code, original action, evidence reviewed, reviewer disposition, and decision timestamp.

Where multiple rules fire, the reviewer should resolve each material condition or document why a downstream action made a later rule unnecessary. The audit trail should still show all codes reported by the engine.

An information request that is satisfied late may require a fresh source pull. Reviewers should check the configured evidence window before treating an uploaded document as current.

Counteroffers should be assessed against the same final evidence package used for the underlying risk decision. A counteroffer is not a way to bypass an unresolved identity, compliance, or fraud blocker.

Notice generation is a separate operational step. Reviewers should confirm that the applicant-facing notice reflects the controlled disposition rather than copying internal shorthand from a case comment.

A reviewer should record the reason an exception was accepted, including any applicable footnote or configured override. 'System allowed it' is not sufficient rationale for audit purposes.

Cases with unresolved source conflicts, unusual fraud patterns, or policy-version ambiguity should remain in Manual Review until the governing configuration and evidence state are identified.

EXHIBIT A — PRESCREEN SELECTION CRITERIA

FIRM OFFER CRITERIA

A Prescreen firm offer is available only after the audience-selection criteria have passed and the record is eligible to move into application intake. The offer does not bypass full underwriting. The purpose of the firm-offer designation in this synthetic manual is to identify a candidate who satisfied a predefined selection state under the marketing population rules. Once the applicant applies, current application data and refreshed verification signals control the final decision.

Firm-offer records should carry the following rule lineage: Table 1 identifies the selection criteria that allowed the audience record to be marketed; Table 3 identifies the prequalification rules applied once the candidate enters the underwriting flow; and Table 4 identifies the refreshed full-underwriting rules applied before final disposition. Codes 101 through 119 represent selected early controls, while 120 through 130 represent additional full-stage controls or reused final-stage decisions. The existence of a firm offer must not be interpreted as evidence that all downstream rules have passed.

Selection criteria should be evaluated with the data snapshot associated with the campaign generation date. If a record becomes ineligible after selection but before application, the application is subject to current underwriting rather than the stale audience state. For operational consistency, the selection file should carry a selection timestamp, campaign identifier, and source-version identifier. Marketing teams should avoid modifying the selected population after export without preserving the original and revised populations so that offer-rate changes can be reconciled.

The following interpretation examples illustrate the intended relationship between exhibits and rule tables. A record blocked by the identity criteria corresponding to Code 101 remains subject to the Blocked SSN rule at underwriting. A record that passed the selection file but now has an active bureau freeze enters the RFAI path under Code 105. A record whose score refreshed below the full-stage cutoff can be declined under Code 121 even though it was not below the prescreen threshold when initially selected. These examples demonstrate why prescreen selection is not a final underwriting approval.

EXHIBIT B — UNDERWRITING DECISION FLOW

New Customer Pre-Qual Process

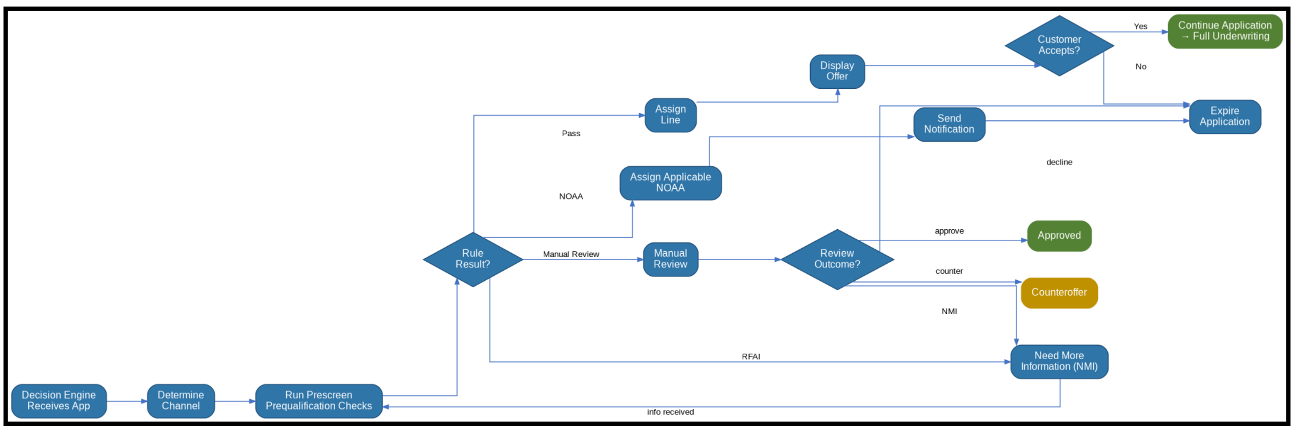

New Customer Full Underwriting Process

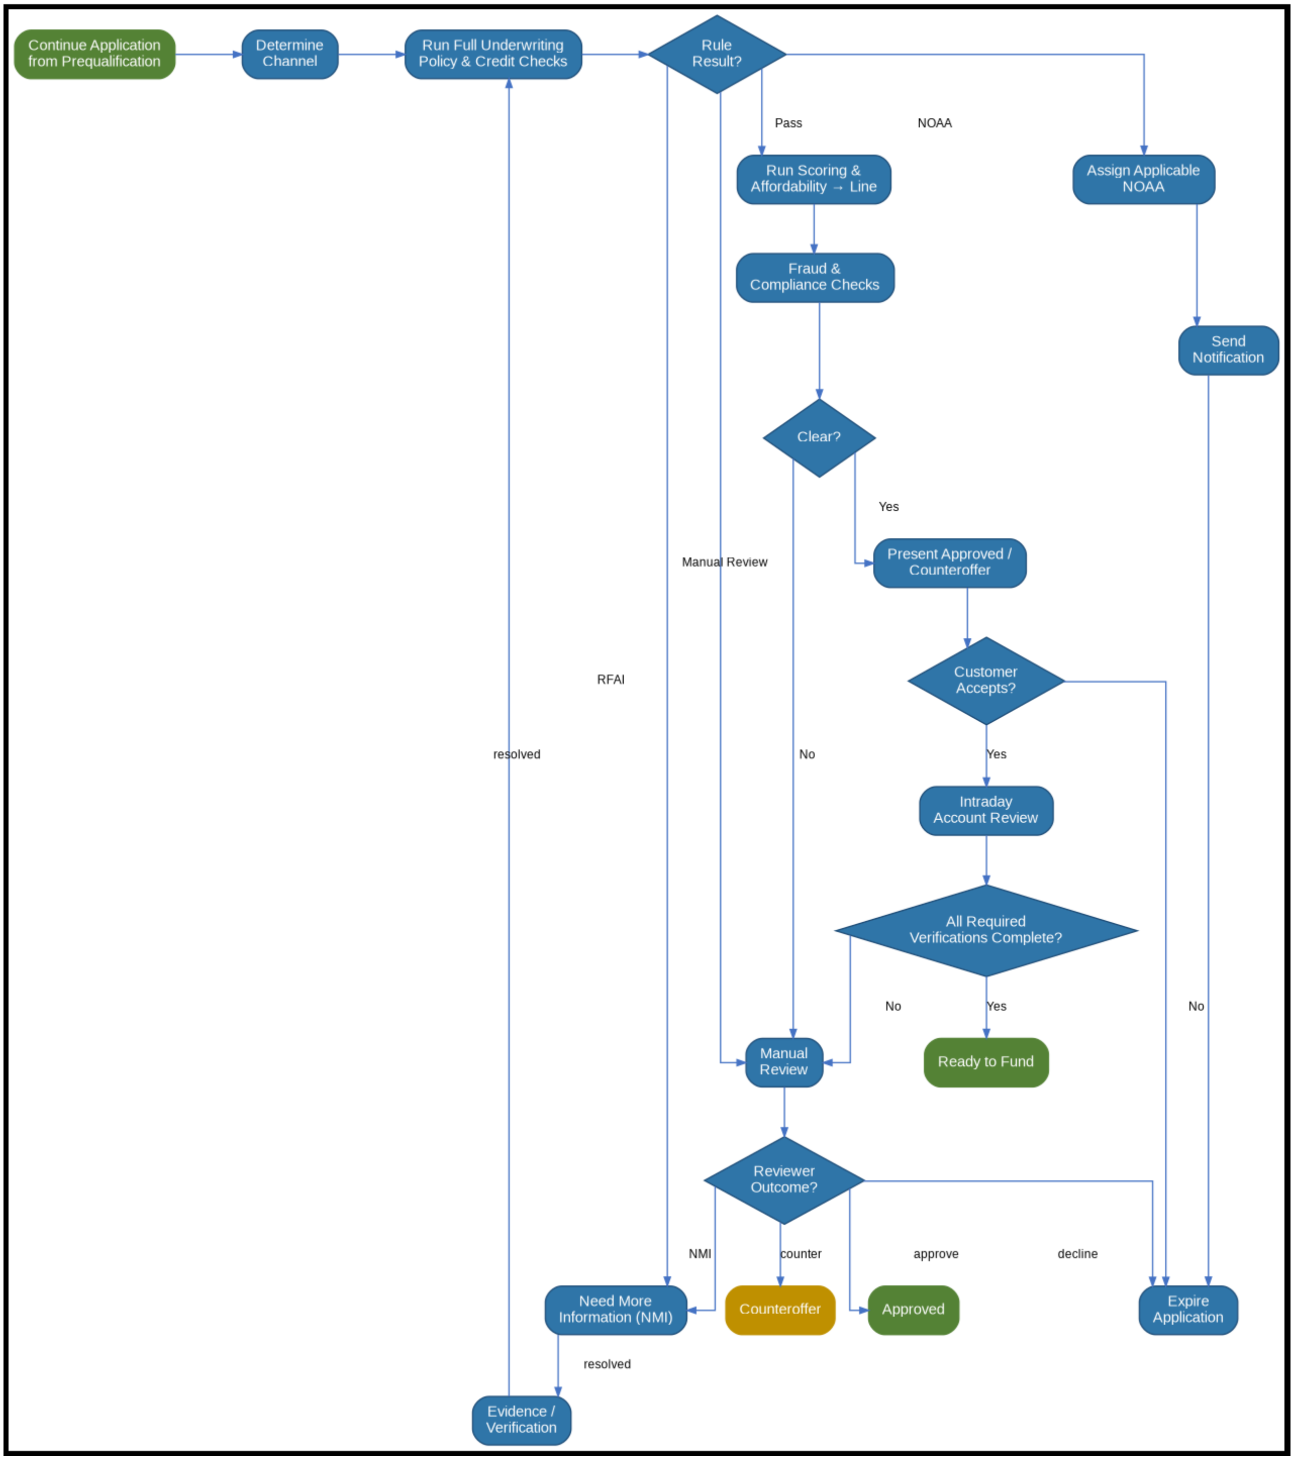

EXHIBIT C — AVAILABLE NOAAs, RFAIs, AND NOIAs

This appendix is the authoritative code dictionary for Tables 3 through 6. Each code is defined with its disposition family and operational meaning. Every code shown in those rule tables appears below, and every appendix code is used at least once in Tables 3 through 6. The two-column format is intentional: the underwriting tables provide the rule-to-code linkage, while Exhibit C provides the code-to-meaning lookup.

Table 9: Available NOAAs

Code	Description
NOAA 101	NOAA — Blocked SSN: SSN matches an active NovaShield identity block. Exception: an administrative duplicate older than 36 months with exact current identity comparison may be treated as passed.
NOAA 102	NOAA — Blocked Phone: phone number matches a high-confidence synthetic fraud suppression record.
NOAA 103	NOAA — Blocked Email: email hash appears on the synthetic high-risk email registry.
NOAA 107	NOAA — Score Floor: representative bureau score is below 566 during prequalification.
NOAA 108	NOAA — Insufficient Income: verified or normalized monthly income is below $1,850.
NOAA 109	NOAA — SSI Threshold: SSI-only income is below the synthetic $1,120 monthly equivalent threshold.
NOAA 110	NOAA — Recent Application Decline: a NovaCred application was declined recently for a non-fraud credit reason; the Prescreen threshold is within 21 days and the full-underwriting overlay is within 10 days after cleared prequalification.
NOAA 111	NOAA — Inquiry Velocity: more than 6 hard inquiries occurred in the prior 120 days, excluding NovaCred soft checks.
NOAA 113	NOAA — OFAC Match: a potential sanctions screening match remains uncleared after automated or secondary comparison.
NOAA 117	NOAA — Existing Exposure: current Flex balance plus the requested amount exceeds the synthetic prequalification exposure cap.
NOAA 118	NOAA — Duplicate Application: identity fingerprint indicates an active duplicate application in the same campaign window.
NOAA 120	NOAA — Blocked Routing: a blocked identity route persists after prequalification and remains present at full underwriting refresh. The defined temporary-network-anomaly exception may override this result when its conditions are satisfied.
NOAA 121	NOAA — Credit Score Confirmation: refreshed representative bureau score is below 580 at full underwriting.
NOAA 123	NOAA — Employment Tenure: employment tenure is under 45 days without an alternate qualifying income source.

Code	Description
NOAA 124	NOAA — Bank Stability: average monthly transaction inflow is below $2,050 over the eligible observation window.
NOAA 125	NOAA — Payment Stress: projected installment-to-net-inflow ratio is above 0.31.
NOAA 126	NOAA — High Inquiry Overlay: more than 9 hard inquiries occurred in 120 days, triggering a stricter final assessment.
NOAA 130	NOAA — Document Integrity: required evidence contains unresolved alteration, corruption, or integrity indicators.
NOAA 133	NOAA — Early-Payoff Pattern: prior NovaCred loan shows three or more rapid payoff-and-reapply cycles in 90 days.
NOAA 135	NOAA — Bank Account Ownership: funding or repayment account ownership cannot be reconciled to an eligible applicant.

Table 10: Available RFAIs

Code	Description
RFAI 105	RFAI — Credit Freeze: bureau security freeze prevents normal file access or remains active at full underwriting refresh. Additional information or a refreshed bureau state is required.
RFAI 106	RFAI — Thin File: bureau file has fewer than two tradelines and no qualifying alternate history; if unresolved after refresh, additional information is required.
RFAI 114	RFAI — FACTA Alert / Resolution: an identity-theft alert requires additional evidence or remains unresolved after document review.
RFAI 115	RFAI — Credit Risk Model Score Vector: the Credit Risk Model score vector required for automated risk ranking is missing, incomplete, or outside the applicable automated scoring process.
RFAI 116	RFAI — Bank Transaction Model Score Vector: the Bank Transaction Model score vector required for automated risk ranking is missing, incomplete, or outside the applicable automated scoring process.
RFAI 122	RFAI — Income Reconciliation: stated/application income differs from verified income by more than 28%, requiring evidence reconciliation.
RFAI 128	RFAI — Offer Age: an offer or campaign artifact is too old at application intake. Prescreen uses an older-than-45-day threshold; Non-Prescreen uses an older-than-30-day threshold.
RFAI 134	RFAI — Channel Attestation: partner-originated application lacks required acquisition attestation.

Code	Description
RFAI 139	RFAI — Address Evidence: address evidence is insufficient to reconcile the current residence after required automated checks.
RFAI 140	RFAI — Funding Method Verification: requested funding method requires additional ownership or destination verification.

Table 11: Available NOIAs

Code	Description
NOIA 136	NOIA — Borrower Notice / Verification Pending: eligible applicant requires a synthetic borrower notice after a documentation event.
NOIA 137	NOIA — Borrower Notice / Data Correction Requested: critical application field was corrected after initial intake and notice issuance.
NOIA 138	NOIA — Final Disposition Pending: final notice generation is pending a downstream policy event.

Table 12: Available Manual Review Codes

Code	Description
Manual Review 104	Manual Review — IP Velocity: application IP exceeds the synthetic rolling 24-hour application threshold; contextual review is required because shared networks or partner traffic can create legitimate velocity.
Manual Review 112	Manual Review — MLA Check / MLA Final: MLA coverage or status remains unresolved at the applicable underwriting stage.
Manual Review 119	Manual Review — Contact Inconsistency: applicant name and contact metadata show an unresolved high-risk mismatch requiring reconciliation.
Manual Review 127	Manual Review — Manual Identity Reconciliation: required core identity fields disagree across two or more sources and require contextual reconciliation.
Manual Review 129	Manual Review — Residence Mismatch: declared/application residence conflicts with validated address information and cannot be resolved automatically.
Manual Review 131	Manual Review — Acquisition Source Risk: Non-Prescreen partner or campaign source has a temporary risk posture above its permitted band.
Manual Review 132	Manual Review — Device Reuse: device fingerprint is linked to 5 or more unrelated active applications during the prior 12 hours.

Table 13: Code Cross-Reference Audit

Code Family	Codes Covered	Rule Tables Referencing the Codes
NOAA	101, 102, 103, 107, 108, 109, 110, 111, 113, 117, 118, 120, 121, 123, 124, 125, 126, 130, 133, 135	Tables 3-6
RFAI	105, 106, 114, 115, 116, 122, 128, 134, 139, 140	Tables 3-6
NOIA	136, 137, 138	Table 6
Manual Review	104, 112, 119, 127, 129, 131, 132	Tables 3-6

Table 14: Fraud Consortium Criteria

Criterion ID	Synthetic Criterion	Severity	Disposition Guidance
FC-01	Device reused across 6+ unrelated applicants in 12 hours	High	Manual Review unless corroborated by another high-confidence signal
FC-02	Contact point reused across 8+ unrelated active applications	High	Manual Review; verify ownership and relationship
FC-03	Identity attributes converge with known synthetic cluster	Critical	NOAA if identity block is active; otherwise Manual Review
FC-04	Application velocity exceeds synthetic consortium burst threshold	Medium	Manual Review; compare network context
FC-05	Repeated bank-account ownership conflicts across applicants	High	NOAA if ownership cannot be resolved
FC-06	Rapid payoff-and-reapply cycle observed across multiple accounts	Medium	Evaluate with Code 133; do not automatically treat one cycle as fraud
FC-07	Shared device used by a household cluster with consistent identity evidence	Low	May be cleared if identity and ownership evidence reconcile
FC-08	Partner traffic cluster shares anomalous network signature	Medium	Evaluate Code 131 and Table 13 before terminal action

Table 14 is deliberately referenced by the Fraud, Verifications, and Manual Review section. The criteria shown here are synthetic pattern definitions used to support retrieval tests in which a reader must follow a forward cross-reference. The consortium criteria do not create new underwriting codes in this manual; where a disposition maps to a rule, the applicable code remains the code in Tables 3 through 6.

Table 15: Override Footnote Register

Footnote	Annotated Rule	Effect
[1]	Blocked SSN / Blocked SSN-style identity block	May be treated as passed when the documented administrative-duplicate exception applies.
[2]	Blocked Routing	May be treated as passed when the defined temporary network anomaly exception and the documented Intraday Account Review criteria are satisfied.
[3]	Scoring	Missing score vector can be treated as passed for continuation when the Null Score condition did not trigger; no score is invented.
[4]	RFAI appendix reference	Points the reviewer to Table 10 for the complete RFAI code list.

APPENDIX OPERATING NOTES

SOURCE PRECEDENCE

When multiple approved sources provide the same attribute, the rule configuration must specify a precedence order. The reviewer should not choose the most favorable source simply because it supports an approval. A source can be considered stale even when technically available; the timestamp and freshness window are part of the data contract.

Source precedence should be encoded as an ordered preference list rather than left to the implementation team's judgment.

When two approved sources disagree, retain both responses and apply the configured tie-break rule; do not overwrite one provider's value with the other.

Fallback use should be monitored because a rise in fallback frequency can change the effective population being evaluated.

EVIDENCE RETENTION

For each terminal or manual result, retain the data element, source, timestamp, rule code, action, and reviewer or system identifier. A screenshot without the underlying source reference is insufficient for reproducibility. Evidence should be retained according to the fictional enterprise retention schedule.

Retention should preserve the evidence artifact, metadata, and evaluation event that consumed it.

Where policy permits replacement or refresh, the prior artifact must remain addressable so the original decision can be reconstructed.

Retention schedules should distinguish applicant-submitted documents from transient provider responses and system-generated evaluations.

CONFIGURATION VERSIONING

The rule engine must record the configuration version that generated the result. A code can retain its meaning while its threshold changes between configuration versions. Monitoring should therefore segment by both rule code and version rather than code alone.

Every production evaluation should record the configuration version that supplied its thresholds, mappings, and rule sequence.

Historical versions must remain readable even after a newer configuration becomes active, because prior decisions may need to be reproduced.

A configuration release should identify changed rules and expected monitoring effects before it is activated.

REPROCESSING

A reprocessed application must carry a new evaluation timestamp and preserve the original result. Reprocessing should be triggered by a defined event such as new bureau data, requested evidence, or a documented technical correction. It should not be used to search repeatedly for a favorable outcome.

Reprocessing should create a new evaluation event rather than mutating the original decision record.

The trigger for reprocessing should be classified, such as source correction, applicant response, or policy-approved retry.

Outputs from a reprocessed event should remain linked to the initiating event so reviewers can compare the before-and-after state.

MONITORING

Rule firing rates should be monitored by underwriting path, acquisition source, month, score band, and application type. Sudden changes may reflect data availability changes rather than true risk movement. Analysts should compare rule firing rate against source completeness and latency.

Monitoring should distinguish changes caused by applicant mix from changes caused by upstream source behavior.

Rule-level metrics should be examined alongside source availability, latency, and fallback rates.

Unexpected shifts should be tied back to the configuration version before a policy conclusion is drawn.

EXCEPTION GOVERNANCE

Exception paths should be rare, documented, and separately reportable. An exception that becomes common is a candidate for policy redesign. Reviewers must not create informal exception categories outside the configured code set.

An exception should identify the rule condition, approving authority, evidence basis, and expiration or scope.

Exceptions must not silently become permanent policy; recurring exceptions should be evaluated for whether the underlying rule needs a formal change.

Reporting should separate approved exceptions from ordinary automated passes so that exception volume remains visible.

DECISION EXPLAINABILITY

The reason presented to an operator should correspond to the rule code and its configured description. User-facing notices may use approved language that differs from the internal rule name, but the mapping should be deterministic and auditable.

Decision explanations should connect the final disposition to the rule or rules that materially affected it.

Applicant-facing language should be generated from controlled mappings rather than exposed internal rule names.

Internal audit detail may retain more technical evidence than the external notice, but both should reconcile to the same decision event.

DATA QUALITY QUARANTINE

If a critical source is structurally invalid, the system should quarantine the application or route it to an appropriate information/manual path. Defaulting nulls to favorable values is prohibited under this fictional policy.

Data quality failures should be isolated from ordinary applicant risk signals when the source value cannot be trusted.

A quarantined record should carry a reason, source, timestamp, and permitted recovery action.

Analysts should monitor quarantine rates by source and path because systematic source defects can resemble changes in underwriting behavior.

SCORE NORMALIZATION

Scores must be normalized to the current scoring schema before entering the line assignment matrix. A score from an incompatible model version cannot be silently mapped by numeric similarity.

Scores from different providers or model versions should be normalized before they are compared with band boundaries.

The normalized value must retain lineage to the raw score and scoring version that produced it.

Band assignment should occur after normalization; a missing or incompatible score should follow the configured missing-score path instead of defaulting to a band.

CROSS-PATH ANALYTICS

Because identical rules reuse codes across paths, enterprise monitoring can compare the rate and downstream outcomes associated with Code 105, Code 106, and other shared controls without building artificial translation tables.

Prescreen and Non-Prescreen results should be compared using path-aware denominators because their applicant populations are constructed differently.

Path comparisons should include acquisition source and configuration version so a channel shift is not mistaken for a rule effect.

Shared rule codes should remain comparable across paths, while path-specific rules should be reported separately.

RULE SEQUENCE AND TERMINALITY

STAGE ENTRY

The underwriting sequence begins only after the application has a valid path identifier and required identity envelope. A record with missing path metadata should not enter a best-guess ruleset. The intent is to make stage entry itself observable and prevent silent routing differences from becoming unexplained policy variance.

For stage entry, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting stage entry should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for stage entry should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for stage entry should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

SHORT-CIRCUIT BEHAVIOR

Terminal actions may short-circuit downstream rules, but the engine should retain the exact point of termination. For example, a high-confidence identity block can stop additional scoring calls, while an RFAI typically permits the application to remain open. The implementation must therefore distinguish terminal decline, temporary hold, and review routing.

For short-circuit behavior, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting short-circuit behavior should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for short-circuit behavior should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for short-circuit behavior should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

MULTI-FIRE CASES

More than one rule can fire in the same stage. The system should preserve all fired codes even when one code determines the immediate action. This allows analysts to understand secondary risk patterns and prevents a later investigation from losing evidence simply because the first terminal rule already ended the workflow.

For multi-fire cases, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting multi-fire cases should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for multi-fire cases should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for multi-fire cases should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

ORDERING CHANGES

Changing rule order can change observable outcomes even if each individual rule is unchanged. Configuration governance should therefore treat ordering as a versioned policy attribute. A month-over-month change in decline mix should be reviewed alongside ordering changes rather than attributed automatically to applicant behavior.

For ordering changes, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting ordering changes should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for ordering changes should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for ordering changes should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

BUREAU DATA HANDLING

REPRESENTATIVE SCORE SELECTION

The synthetic bureau layer may return more than one score or no score. The configured representative-score selection method must be deterministic and recorded. A reviewer should not select a score based on which score produces the better outcome. If the expected score source is absent, the application follows the documented missing-data behavior rather than a favorable default.

Source lineage for representative score selection should retain the raw bureau response together with the normalized field consumed by the rule, including the provider timestamp.

Schema changes affecting representative score selection should be visible in the mapping layer. An incompatible field should produce a controlled missing-data state rather than a synthetic value.

When multiple bureau responses are available for representative score selection, the decision record should identify which response was selected and why it was considered current.

Monitoring for representative score selection should distinguish source-quality failures from genuine applicant conditions, using response completeness and refresh timing as separate diagnostic measures.

FREEZE PROCESSING

A freeze is treated as an information state in the early rules because the applicant may be able to resolve it or provide another approved evidence path. The full-underwriting refresh repeats the condition so that a stale prequalification result does not accidentally survive into final decisioning.

Source lineage for freeze processing should retain the raw bureau response together with the normalized field consumed by the rule, including the provider timestamp.

Schema changes affecting freeze processing should be visible in the mapping layer. An incompatible field should produce a controlled missing-data state rather than a synthetic value.

When multiple bureau responses are available for freeze processing, the decision record should identify which response was selected and why it was considered current.

Monitoring for freeze processing should distinguish source-quality failures from genuine applicant conditions, using response completeness and refresh timing as separate diagnostic measures.

THIN-FILE PROCESSING

Thin file is not a synonym for poor credit. The rule is specifically about insufficient observable history for the fictional automated strategy. An applicant can have no adverse tradelines and still trigger the thin-file rule because the model does not have enough evidence under the defined control.

Source lineage for thin-file processing should retain the raw bureau response together with the normalized field consumed by the rule, including the provider timestamp.

Schema changes affecting thin-file processing should be visible in the mapping layer. An incompatible field should produce a controlled missing-data state rather than a synthetic value.

When multiple bureau responses are available for thin-file processing, the decision record should identify which response was selected and why it was considered current.

Monitoring for thin-file processing should distinguish source-quality failures from genuine applicant conditions, using response completeness and refresh timing as separate diagnostic measures.

INQUIRY HANDLING

Inquiry count uses a defined observation window and exclusions. Soft NovaCred checks are excluded from the count in the early rule. The stricter full-stage overlay uses a higher threshold and remains a separate code so monitoring can distinguish normal prequalification pressure from a final-stage overlay.

Source lineage for inquiry handling should retain the raw bureau response together with the normalized field consumed by the rule, including the provider timestamp.

Schema changes affecting inquiry handling should be visible in the mapping layer. An incompatible field should produce a controlled missing-data state rather than a synthetic value.

When multiple bureau responses are available for inquiry handling, the decision record should identify which response was selected and why it was considered current.

Monitoring for inquiry handling should distinguish source-quality failures from genuine applicant conditions, using response completeness and refresh timing as separate diagnostic measures.

IDENTITY, CONTACT, AND NETWORK CONTROLS

IDENTITY BLOCKS

A block record is treated as a high-confidence control when its status is active and its identity linkage is exact enough for the configured rule. Footnote [1] is intentionally narrow. It permits an exception for a resolved administrative duplicate but does not authorize a general reviewer override whenever the reviewer disagrees with the block.

For identity blocks, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting identity blocks should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for identity blocks should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for identity blocks should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

PHONE CONTROLS

Phone suppression can be generated by synthetic fraud signals or repeated reuse patterns. The implementation should preserve the reason category and confidence. A reused phone is not automatically equivalent to a blocked phone; the policy separates the high-confidence blocked condition from contextual velocity signals.

For phone controls, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting phone controls should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for phone controls should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for phone controls should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

EMAIL CONTROLS

Email hash matching is used to avoid storing raw contact values in the synthetic risk registry. A blocked hash can trigger a NOAA, while unresolved name/contact mismatch is routed to Manual Review. This distinction creates a retrieval edge case in which two contact-related rules have different actions and codes.

For email controls, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting email controls should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for email controls should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for email controls should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

IP CONTROLS

IP Velocity is a contextual trigger. Shared networks, partner environments, and household behavior can create legitimate reuse. That is why Code 104 is Manual Review rather than an automatic decline. Table 13 and Table 14 provide additional context for interpreting the signal.

For ip controls, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting ip controls should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for ip controls should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for ip controls should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

DEVICE CONTROLS

Device Reuse is unique to the Non-Prescreen path in this manual. It detects a pattern across multiple active applications but does not establish fraud by itself. Reviewers should compare the device signal against identity consistency, applicant relationship, and consortium criteria before choosing the final disposition.

For device controls, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting device controls should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for device controls should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for device controls should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

INCOME AND AFFORDABILITY CALCULATION STANDARDS

MONTHLY NORMALIZATION

Income values should be normalized to a monthly amount using the source-specific convention. Salary, weekly pay, annual contract income, and business inflow may use different conversion logic. The normalized amount should be stored alongside the original amount and source frequency so that a reviewer can reproduce the calculation.

For monthly normalization, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting monthly normalization should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for monthly normalization should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for monthly normalization should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

DUPLICATE INCOME STREAMS

The same employment income should not be counted once from payroll and again from bank inflows. When two sources represent the same underlying stream, the controlling source should be chosen using the tier rules. The other source can serve as corroboration but not as additive income unless the policy explicitly identifies a separate stream.

For duplicate income streams, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting duplicate income streams should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for duplicate income streams should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for duplicate income streams should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

SELF-EMPLOYED INCOME

Tier 3 is intentionally different for self-employed applicants. Gross business receipts can overstate spendable repayment capacity. The synthetic policy therefore requires eligible expense adjustments and a reconciliation between business records and bank activity before monthly income is finalized.

For self-employed income, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting self-employed income should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for self-employed income should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for self-employed income should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

PAYMENT STRESS

Payment Stress is a final-stage affordability control and uses net inflow rather than a generic score. Analysts should inspect both numerator and denominator if the rule fires unexpectedly. A sudden increase in firing rate can arise from changes in transaction categorization rather than true deterioration in applicants.

For payment stress, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting payment stress should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for payment stress should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for payment stress should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

COUNTEROFFER LOGIC

An affordability-based counteroffer should preserve the original requested amount and the adjusted amount. The system should not overwrite the original request, because later monitoring may need to distinguish demand from capacity. A counteroffer is a new proposed line, not evidence that the original request was approved.

For counteroffer logic, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting counteroffer logic should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for counteroffer logic should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for counteroffer logic should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

COMPLIANCE CONTROL OPERATIONS

MLA STATUS

The synthetic MLA check is intentionally a Manual Review state when status cannot be resolved automatically. It is not treated as a simple score. The reviewer must establish which status was returned, when it was obtained, and whether the ambiguity can be resolved with an approved source.

For mla status, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting mla status should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for mla status should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for mla status should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

OFAC SCREENING

The synthetic OFAC rule treats an uncleared match as a NOAA. Secondary comparison can clear some potential matches, but the final result must be recorded. An automated mismatch should not be hidden simply because the applicant otherwise passes credit and income checks.

For ofac screening, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting ofac screening should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for ofac screening should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for ofac screening should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

FACTA ALERT

A FACTA alert is treated as an RFAI because additional identity evidence may resolve the condition. The alert can appear both at prequalification and final refresh. The reused code means downstream monitoring can compare the control across stages.

For facta alert, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting facta alert should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for facta alert should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for facta alert should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

NOTICE GENERATION

NOIA codes 136-138 are document-generation events, not substitutes for credit outcomes. They should be generated in addition to the relevant underwriting decision record when the configured notice condition occurs. This makes the appendix deliberately contain a third family of codes that cannot be mistaken for NOAA or RFAI.

For notice generation, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting notice generation should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for notice generation should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for notice generation should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

SCORING AND MODEL OUTPUT HANDLING

SCORE LINEAGE

The scoring rule references two synthetic models: the Credit Risk Model and the Bank Transaction Model. The rule intentionally uses multiple codes so that model-output failures can be diagnosed by component. If both outputs are available but combined score is weak, the scoring row is still evaluated through the configured action; if one output is missing, the specific code associated with that component can identify the failure source.

For score lineage, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting score lineage should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for score lineage should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for score lineage should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

MISSING SCORES

Footnote [3] provides a narrow continuation override. It does not manufacture a score and it does not guarantee approval. The downstream workflow may still require another evidence path before line assignment, particularly when the score bands cannot be populated safely.

For missing scores, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting missing scores should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for missing scores should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for missing scores should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

SCORE CUT BOUNDARIES

Band boundaries are inclusive or exclusive exactly as shown in the matrix labels. Implementations must not infer boundaries from neighboring cells. Boundary tests should include values immediately below, at, and immediately above each cut so that drift caused by comparison operators is caught.

For score cut boundaries, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting score cut boundaries should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for score cut boundaries should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for score cut boundaries should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

LINE FLOOR AND CAP

The minimum floor and maximum cap are independent controls. A positive matrix adjustment cannot push the line above the cap, and a negative adjustment cannot reduce the line below the floor. These controls are applied after the matrix cell is selected.

For line floor and cap, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting line floor and cap should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for line floor and cap should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for line floor and cap should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

MONITORING SCORE MIX

Monitoring should examine the distribution of score bands as well as the frequency of each rule. A stable rule-firing rate with a shifting score mix can still materially change the final line distribution.

For monitoring score mix, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting monitoring score mix should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for monitoring score mix should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for monitoring score mix should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

APPLICATION REWORK, NMI, AND REMEDIATION

RFAI LIFECYCLE

An RFAI starts an information-request lifecycle. The request should specify what evidence is needed, why it is needed, when it expires, and what rule code created it. When evidence arrives, the application is re-evaluated against the current stage. The old RFAI event remains in history even if the condition is resolved.

For rfai lifecycle, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting rfai lifecycle should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for rfai lifecycle should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for rfai lifecycle should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

MULTIPLE REQUESTS

If several RFAI codes fire, operations should consolidate requests where the same document can satisfy multiple rules, but the audit record must preserve the relationship from each request back to its source code. This prevents a single uploaded document from obscuring which policy conditions were actually resolved.

For multiple requests, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting multiple requests should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for multiple requests should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for multiple requests should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

EXPIRED EVIDENCE

Evidence received after its validity window should not be treated as automatically current. The case may need a refreshed bureau or income pull. This is a timing issue, not necessarily an applicant risk issue.

For expired evidence, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting expired evidence should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for expired evidence should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for expired evidence should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.

REOPEN CONTROLS

A declined application should not be manually reopened simply because an applicant calls. Reopening should use a configured event type and preserve the original decline. A new application may be appropriate when the original evidence is materially stale.

For reopen controls, the evidence record should retain the source value, normalized interpretation, rule context, and timestamp needed to explain the result. The control is evaluated from the approved inputs available at that stage.

A change affecting reopen controls should create an observable event rather than silently altering an earlier decision. Re-evaluation should state what changed, which rule was rerun, and whether any prior action remains in force.

Exception handling for reopen controls should be explicit. When a fallback, override, or manual disposition is permitted, the record should identify the condition that allowed it and the authority or configuration that governed the exception.

Monitoring for reopen controls should look beyond raw firing counts. Reviewers should compare the mix of source outcomes, path, configuration version, and downstream disposition so operational drift is not mistaken for a policy change.